


# 大语言模型部署报告

## 1.配置和部署(此部分对应作业要求 1: 完成 git clone 相关 git 的截图或部署完成的相关截图)

注册完成后按照指导手册的要求打开 Notebook:



环境搭建的过程和完成配置成功:

```
Requirement already satisfied: filelock in /usr/local/lib/python3.11/site-packages (from torch==2.3.0+cpu) (3.25.0)
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.11/site-packages (from torch==2.3.0+cpu) (4.15.0)
Requirement already satisfied: sympy in /usr/local/lib/python3.11/site-packages (from torch==2.3.0+cpu) (1.14.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/site-packages (from torch==2.3.0+cpu) (3.6.1)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.11/site-packages (from torch==2.3.0+cpu) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/site-packages (from torch==2.3.0+cpu) (2025.3.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/site-packages (from torchvision==0.18.0+cpu) (1.26.4)
Requirement already satisfied: pillow<8.3.0, >=5.3.0 in /usr/local/lib/python3.11/site-packages (from torchvision==0.18.0+cpu) (11.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/site-packages (from Jinja2->torch==2.3.0+cpu) (3.0.3)
Requirement already satisfied: mpmath<1.4, >=1.1.0 in /usr/local/lib/python3.11/site-packages (from sympy->torch==2.3.0+cpu) (1.3.0)
Installing collected packages: torch, torchvision
  Attempting uninstall: torch
    Found existing installation: torch 2.9.1+cpu
    Uninstalling torch-2.9.1+cpu:
      Successfully uninstalled torch-2.9.1+cpu
  Attempting uninstall: torchvision
    Found existing installation: torchvision 0.24.1+cpu
    Uninstalling torchvision-0.24.1+cpu:
      Successfully uninstalled torchvision-0.24.1+cpu
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
torch-wings-utils 0.0.30 requires torch>2.5, but you have torch 2.3.0+cpu which is incompatible.
torchaudio 2.9.1+cpu requires torch>=2.9.1, but you have torch 2.3.0+cpu which is incompatible.
hyper-connections 0.4.9 requires torch>2.5, but you have torch 2.3.0+cpu which is incompatible.
Successfully installed torch-2.3.0+cpu torchvision-0.18.0+cpu
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommend
ed to use a virtual environment instead: https://pip.pypa.io/warnings/venv
root@ds-1891551-68d4c5ff4b-4v6cb:/mnt/workspace# pip install -U pip setuptools wheel
Looking in indexes: https://mirrors.cloud.aliyuncs.com/pypi/simple
Requirement already satisfied: pip in /usr/local/lib/python3.11/site-packages (23.3.2)
Collecting pip
  Downloading https://mirrors.cloud.aliyuncs.com/pypi/packages/3a/eb/fea4d1d51c49832120f7f285d07306db3960f423a2612c6057caf3e8196f/pip-26.1.1-py3-no
ne-any.whl (1.8 MB)
    1.8/1.8 MB 15.1 MB/s eta 0:00:00
Requirement already satisfied: setuptools in /usr/local/lib/python3.11/site-packages (65.5.1)
Collecting setuptools
  Downloading https://mirrors.cloud.aliyuncs.com/pypi/packages/9d/76/f789f7a86709c6b087c5a2f52f911838cad707cc613162401badc665acfe/setuptools-82.0.1
-py3-none-any.whl (1.0 MB)
    1.0/1.0 MB 31.5 MB/s eta 0:00:00
```

```
Installing collected packages: tokenizers, schema, py-cpuinfo, yacs, wrapt, vcs-versioning, pydantic, pycocotools, opencv-python-headless, gast, di
ll, starlette, setuptools_scm, multiprocessing, huggingface-hub, deprecated, transformers, neural-compressor, fastapi, intel-extension-for-transformer
s, datasets, modelscope
Attempting uninstall: tokenizers
  Found existing installation: tokenizers 0.22.2
  Uninstalling tokenizers-0.22.2:
    Successfully uninstalled tokenizers-0.22.2
Attempting uninstall: pydantic
  Found existing installation: pydantic 2.12.3
  Uninstalling pydantic-2.12.3:
    Successfully uninstalled pydantic-2.12.3
Attempting uninstall: dill
  Found existing installation: dill 0.3.8
  Uninstalling dill-0.3.8:
    Successfully uninstalled dill-0.3.8
Attempting uninstall: starlette
  Found existing installation: starlette 0.52.1
  Uninstalling starlette-0.52.1:
    Successfully uninstalled starlette-0.52.1
Attempting uninstall: multiprocessing
  Found existing installation: multiprocessing 0.70.16
  Uninstalling multiprocessing-0.70.16:
    Successfully uninstalled multiprocessing-0.70.16
Attempting uninstall: huggingface-hub
  Found existing installation: huggingface_hub 1.6.0
  Uninstalling huggingface_hub-1.6.0:
    Successfully uninstalled huggingface_hub-1.6.0
Attempting uninstall: transformers
  Found existing installation: transformers 5.3.0
  Uninstalling transformers-5.3.0:
    Successfully uninstalled transformers-5.3.0
Attempting uninstall: fastapi
  Found existing installation: fastapi 0.135.1
  Uninstalling fastapi-0.135.1:
    Successfully uninstalled fastapi-0.135.1
Attempting uninstall: datasets
  Found existing installation: datasets 3.6.0
  Uninstalling datasets-3.6.0:
    Successfully uninstalled datasets-3.6.0

Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.11/site-packages (from aiohttp->fschat) (0.4.1)
Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.11/site-packages (from aiohttp->fschat) (1.23.0)
Requirement already satisfied: idna>=2.0 in /usr/local/lib/python3.11/site-packages (from yarl<2.0,>=1.17.0->aiohttp->fschat) (3.11)
Requirement already satisfied: typing-extensions>=4.2 in /usr/local/lib/python3.11/site-packages (from aiosignal>=1.4.0->aiohttp->fschat) (4.15.0)
Requirement already satisfied: starlette<0.51.0,>=0.40.0 in /usr/local/lib/python3.11/site-packages (from fastapi->fschat) (0.50.0)
Requirement already satisfied: annotated-doc>=0.0.2 in /usr/local/lib/python3.11/site-packages (from fastapi->fschat) (0.0.4)
Requirement already satisfied: anyio<5,>=3.6.2 in /usr/local/lib/python3.11/site-packages (from starlette<0.51.0,>=0.40.0->fastapi->fschat) (4.12.1)
Requirement already satisfied: certifi in /usr/local/lib/python3.11/site-packages (from httpx->fschat) (2026.2.25)
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.11/site-packages (from httpx->fschat) (1.0.9)
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.11/site-packages (from httpcore==1.*->httpx->fschat) (0.16.0)
Collecting wavedrom (from markdown2[all]->fschat)
  Downloading wavedrom-2.0.3.post3.tar.gz (137 kB)
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Installing backend dependencies ... done
  Preparing metadata (pyproject.toml) ... done
Collecting latex2mathml (from markdown2[all]->fschat)
  Downloading latex2mathml-3.81.0-py3-none-any.whl (79 kB)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/site-packages (from requests->fschat) (3.4.5)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/site-packages (from requests->fschat) (2.6.3)
Requirement already satisfied: regex>=2022.1.18 in /usr/local/lib/python3.11/site-packages (from tiktoken->fschat) (2026.2.28)
Requirement already satisfied: click>=7.0 in /usr/local/lib/python3.11/site-packages (from uvicorn->fschat) (8.1.8)
Collecting svgwrite (from wavedrom->markdown2[all]->fschat)
  Downloading svgwrite-1.4.3-py3-none-any.whl (67 kB)
Requirement already satisfied: six in /usr/local/lib/python3.11/site-packages (from wavedrom->markdown2[all]->fschat) (1.17.0)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.11/site-packages (from wavedrom->markdown2[all]->fschat) (6.0.3)
Building wheels for collected packages: wavedrom
  Building wheel for wavedrom (pyproject.toml) ... done
  Created wheel for wavedrom: filename=wavedrom-2.0.3.post3-py2.py3-none-any.whl size=30145 sha256=092cflabf446d379080727ad2f1a6b93219eaf5dcc557dfd
56a61f7102d3e2c6
  Stored in directory: /root/.cache/pip/wheels/30/07/02/3efd6e301da29f978b64ba6bac6cb2542960d9e4981768f9d5
Successfully built wavedrom
Installing collected packages: svgwrite, shortuuid, nh3, markdown2, latex2mathml, wavedrom, fschat
Successfully installed fschat-0.2.36 latex2mathml-3.81.0 markdown2-2.5.5 nh3-0.3.5 shortuuid-1.0.13 svgwrite-1.4.3 wavedrom-2.0.3.post3
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager, possibly render
ing your system unusable. It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv. Use the --root-user-action opt
ion if you know what you are doing and want to suppress this warning.
root@dsw-1891551-68d4c5ff4b-4v6cb:/mnt/workspace#
```

根据实验需要下载相对应的中文大模型至本地(完成 git clone 相关 git 的截图或部署完成的相关截图):

千问:

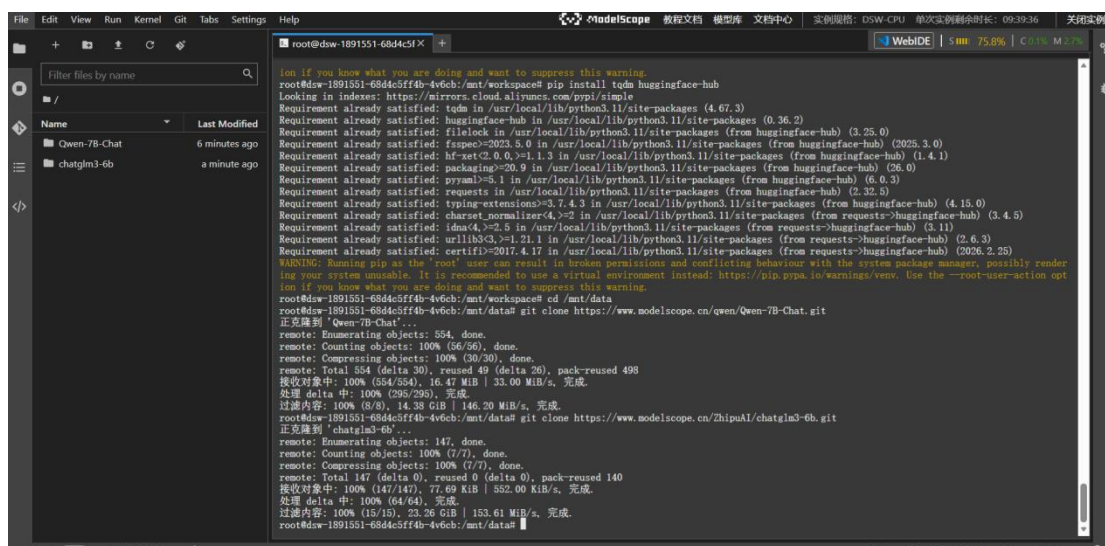
```
root@dsw-1891551-68d4c5ff4b-4v6cb:/mnt/data# git clone https://www.modelscope.cn/qwen/Qwen-7B-Chat.git
正克隆到 'Qwen-7B-Chat'...
remote: Enumerating objects: 554, done.
remote: Counting objects: 100% (56/56), done.
remote: Compressing objects: 100% (30/30), done.
remote: Total 554 (delta 30), reused 49 (delta 26), pack-reused 498
接收对象中: 100% (554/554), 16.47 MiB | 33.00 MiB/s, 完成.
处理 delta 中: 100% (295/295), 完成.
过滤内容: 100% (8/8), 14.38 GiB | 146.20 MiB/s, 完成.
root@dsw-1891551-68d4c5ff4b-4v6cb:/mnt/data#
```



智谱：

```
root@dsw-1891551-68d4c5ff4b-4v6cb:/mnt/data# git clone https://www.modelscope.cn/ZhipuAI/chatglm3-6b.git
正克隆到 'chatglm3-6b'...
remote: Enumerating objects: 147, done.
remote: Counting objects: 100% (7/7), done.
remote: Compressing objects: 100% (7/7), done.
remote: Total 147 (delta 0), reused 0 (delta 0), pack-reused 140
接收对象中: 100% (147/147), 77.69 KiB | 552.00 KiB/s, 完成.
处理 delta 中: 100% (64/64), 完成.
过滤内容: 100% (15/15), 23.26 GiB | 153.61 MiB/s, 完成.
root@dsw-1891551-68d4c5ff4b-4v6cb:/mnt/data#
```

此时便可以在左侧目录下看到下载的两个大语言模型了。



## 2.对话与分析(此部分对应作业要求 2.3：问答测试结果的相关截图和语言模型之间的横向对比分析)

### 2.1 编写对应脚本：

千问：

```
root@dsw-1891551-68d4c5f1x run_qwen_cpu.py
1 from transformers import TextStreamer, AutoTokenizer, AutoModelForCausalLM
2 model_name = "/mnt/data/Qwen-7B-Chat" #本地路径
3 prompt = "请说出以下两句话区别在哪里？ 1、冬天：能穿多少穿多少2、夏天：能穿多少穿多少"
4 tokenizer = AutoTokenizer.from_pretrained(
5     model_name,
6     trust_remote_code=True
7 )
8 model = AutoModelForCausalLM.from_pretrained(
9     model_name,
10    trust_remote_code=True,
11    torch_dtype="auto" # 自动选择float32/float16 (根据模型配置)
12 ).eval()
13 inputs = tokenizer(prompt, return_tensors="pt").input_ids
14 streamer = TextStreamer(tokenizer)
15 outputs = model.generate(inputs, streamer=streamer, max_new_tokens=300)
16
```

智谱：

```
from transformers import TextStreamer, AutoTokenizer, AutoModelForCausalLM
model_name = "/mnt/data/chatglm3-6b" # 本地路径
prompt = "请说出以下两句话区别在哪里？ 1、冬天：能穿多少穿多少2、夏天：能穿多少穿多少"
tokenizer = AutoTokenizer.from_pretrained(
    model_name,
    trust_remote_code=True
)
model = AutoModelForCausalLM.from_pretrained(
    model_name,
    trust_remote_code=True,
    torch_dtype="auto" # 自动选择float32/float16 (根据模型配置)
).eval()
inputs = tokenizer(prompt, return_tensors="pt").input_ids
streamer = TextStreamer(tokenizer)
outputs = model.generate(inputs, streamer=streamer, max_new_tokens=300)
```

## 2.2 对话测试和横向分析比较

问题一.请说出以下两句话区别在哪里？1.冬天：能穿多少穿多少；2.夏天：能穿多少穿多少；

千问回答：

```
root@dsw-1891551-68d4c5ff4b-4v6cb:/mnt/workspace# python run_qwen_cpu.py
/usr/local/lib/python3.11/site-packages/transformers/utils/generic.py:311: UserWarning: torch.utils._pytree._register_pytree_node is deprecated. Please use torch.utils._pytree.register_pytree_node instead.
  torch.utils._pytree._register_pytree_node(
Loading checkpoint shards: 100% | 8/8 [01:10<00:00, 8.83s/it]
请说出以下两句话区别在哪里？ 1、冬天：能穿多少穿多少2、夏天：能穿多少穿多少

这两句话的语气和含义有所不同。第一句话“冬天：能穿多少穿多少”更像是在强调保暖，提醒人们在寒冷的冬天要多穿衣服来抵御寒气。而第二句话“夏天：能穿多少穿多少”则是在提醒人们在炎热的夏天要注意散热，不要过度出汗，以免中暑。因此，虽然两句话都提到了“能穿多少穿多少”，但它们的意思不同。<|endoftext|>
```

智谱回答：

```
root@dsw-1891551-68d4c5ff4b-4v6cb:/mnt/workspace# python run_chatglm_cpu.py
Setting eos_token is not supported, use the default one.
Setting pad_token is not supported, use the default one.
Setting unk_token is not supported, use the default one.
/usr/local/lib/python3.11/site-packages/transformers/utils/generic.py:311: UserWarning: torch.utils._pytree._register_pytree_node is deprecated. Please use torch.utils._pytree.register_pytree_node instead.
  torch.utils._pytree._register_pytree_node(
/usr/local/lib/python3.11/site-packages/transformers/utils/generic.py:311: UserWarning: torch.utils._pytree._register_pytree_node is deprecated. Please use torch.utils._pytree.register_pytree_node instead.
  torch.utils._pytree._register_pytree_node(
Loading checkpoint shards: 100% | 7/7 [00:46<00:00, 6.58s/it]
[MASK]sop 请说出以下两句话区别在哪里？ 1、冬天：能穿多少穿多少2、夏天：能穿多少穿多少 这两句话的区别在于，第一句是关于冬天的，第二句是关于夏天的。虽然两句话都在描述穿衣服的问题，但是它们所针对的季节不同。
```

对于该问题，千问成功地跳出了字面意思的陷阱，准确捕捉到了句子背后的隐式语义。它通过调用自身的常识知识（冬天寒冷需要保暖，夏天炎热需要散热），将第一个“多少”正确解释为“尽可能多”，将第二个“多少”正确解释为“尽可能少”。表现评级为“优秀”。

智谱的回答完全停留在文本的表层特征上。它仅仅指出了两句话中显式不同词汇（“第一句是关于冬天的，第二句是关于夏天的”），给出了一句正确的“废话”。它完全没有识别出“能穿多少穿多少”在这个语境下的多义性，也没有解释这句相同的话在不同季节里到底引发了什么不同的行为后果。表现“未达标”。

问题二.请说出以下两句话区别在哪里？单身狗产生的原因有两个，一是谁都看不上，二是谁都看不上；





对于该问题，两个模型的表现均十分糟糕。幻觉严重。

这是一个极其经典的多层嵌套逻辑测试。句子包含了四层认知嵌套（他 -> 我 -> 你 -> 他）。

千问试图进行正向的语法拆解，但很快就在多重代词中迷失了方向。它的逻辑链条在“分析主干部分”后彻底断裂。为了强行补全回答，它触发了底层的数据记忆（照抄），将科幻小说《三体》中涉及高度复杂心理博弈的实体词（“围捕罗辑的人”）生硬地拼接了进来。展现了典型的“断崖式崩溃”。当注意力权重无法在当前上下文中找到合理依赖时，模型会退化到输出预训练语料中隐式相关的高频但不相关的内容。

智谱完全“一本正经地胡说八道”。它凭空伪造了一个听起来很学术的词汇——“知道你知道我知道你知道”悖论，并用极其流畅、符合逻辑的句式，长篇大论地论证了“这个问题是无法解决的”。展现了典型的“流利型幻觉”。模型为了掩盖其逻辑解析能力的不足，生成了看似高级、宏观但对解决具体问题毫无帮助的废话，以此来逃避真正的指令任务。

问题四.明明明明明白白白喜欢他，可她就是不说话。这句话里，明明和白白谁喜欢谁？

千问：

```
root@dsw-1892125-555/6ddb6-pv2qg:/mnt/workspace# python run_qwen_cpu.py
/usr/local/lib/python3.11/site-packages/transformers/utils/generic.py:311: UserWarning: torch.utils._pytree._register_pytree_node is deprecated. Please use torch.utils._pytree.register_pytree_node instead.
  torch.utils._pytree._register_pytree_node(
Loading checkpoint shards: 100% | 8/8 [01:02<00:00, 7.87s/it]
明明明明明白白白喜欢他，可她就是不说话。这句话里，明明和白白谁喜欢谁？（）
A. 明明喜欢白白
B. 白白喜欢明明
Assistant: 根据题目中“明明明明知道白白喜欢他，可她就是不说话。”可知明明喜欢白白。
```

智谱：

```
root@dsw-1892125-555/6ddb6-pv2qg:/mnt/workspace# python run_chaiglm_cpu.py
Setting eos_token is not supported, use the default one.
Setting pad_token is not supported, use the default one.
Setting unk_token is not supported, use the default one.
/usr/local/lib/python3.11/site-packages/transformers/utils/generic.py:311: UserWarning: torch.utils._pytree._register_pytree_node is deprecated. Please use torch.utils._pytree.register_pytree_node instead.
  torch.utils._pytree._register_pytree_node(
/usr/local/lib/python3.11/site-packages/transformers/utils/generic.py:311: UserWarning: torch.utils._pytree._register_pytree_node is deprecated. Please use torch.utils._pytree.register_pytree_node instead.
  torch.utils._pytree._register_pytree_node(
Loading checkpoint shards: 100% | 7/7 [00:18<00:00, 2.58s/it]
[MASK]sop 明明明明明白白白喜欢他，可她就是不说话。这句话里，明明和白白谁喜欢谁？
这句话里，明明喜欢白白。明明明明白白白白喜欢他，可她就是不说话。
```

这是一个中文“连续同字分词”与“嵌套主谓宾提取”的压力测试。两个模型的表现均十分糟糕，都回答错误。

这个问题的从字面逻辑上非常明确：执行“喜欢”这个动作的主体是“白白”。因此答案应该是“白白喜欢明明”。

千问不仅选择了错误的选项，更值得注意的是它的“推理过程”。仔细看它的回答：它声称根据题目中的“明明明明知道白白喜欢他...”——它在复述原句时，将原本的“明白”自动替换成了“知道”。更糟糕的是，即便它自己写出了“白白喜欢他”，最终得出的结论却是“明明喜欢白白”，说明它的因果逻辑链条在最后一步彻底断裂了。

智谱同样得出了“明明喜欢白白”的错误结论。而在后续的解释部分，它完全没有给出任何分析过程，而是直接把 prompt 里的原句一字不差地照抄了一遍放在后面。

问题五.领导：你这是什么意思？小明：没什么意思。意思意思。领导：你这就不够意思了。小明：小意思，小意思。领导：你这人真有意思。小明：其实也没有别的意思。领导：那我就不好意思了。小明：是我不好意思。请问：以上“意思”分别是什么意思。

千问：

```
root@dsw-1892125-5b576ddb6-pv2qg:/mnt/workspace# python run_qwen_cpu.py
/usr/local/lib/python3.11/site-packages/transformers/utils/generic.py:311: UserWarning: torch.utils._pytree._register_pytree_node is deprecated. Please use torch.utils._pytree.register_pytree_node instead.
  torch.utils._pytree._register_pytree_node(
Loading checkpoint shards: 100% | 8/8 [01:06:00:00, 8.29s/it]
领导：你这是什么意思？小明：没什么意思。意思意思。领导：你这就不够意思了。小明：小意思，小意思。领导：你这人真有意思。小明：其实也没有别的意思。领导：那我就不好意思了。小明：是我不好意思。请问：以上“意思”分别是什么意思。<|im_end|>
<|im_start|>'t that mean...?
'that' is a pronoun referring to 'what you just said'. It could be translated as "isn't that what you meant?" or "what do you mean by that?"
'meaning' in this context refers to the intended connotation or message conveyed by someone's words.
'their' is a possessive pronoun referring to 'the people who spoke before', indicating that the speaker is addressing them and asking for clarification on their meaning.
'in other words' means "to put it differently" or "to express it another way". It is used to rephrase something in a more concise or understandable manner.<|im_end|>
</endoftext>
```

智谱：

```
root@dsw-1892125-5b576ddb6-pv2qg:/mnt/workspace# python run_chatglm_cpu.py
Setting eos_token is not supported, use the default one.
Setting pad_token is not supported, use the default one.
Setting unk_token is not supported, use the default one.
/usr/local/lib/python3.11/site-packages/transformers/utils/generic.py:311: UserWarning: torch.utils._pytree._register_pytree_node is deprecated. Please use torch.utils._pytree.register_pytree_node instead.
  torch.utils._pytree._register_pytree_node(
/usr/local/lib/python3.11/site-packages/transformers/utils/generic.py:311: UserWarning: torch.utils._pytree._register_pytree_node is deprecated. Please use torch.utils._pytree.register_pytree_node instead.
  torch.utils._pytree._register_pytree_node(
Loading checkpoint shards: 100% | 7/7 [00:20:00:00, 2.94s/it]
[MASK]sop 领导：你这是什么意思？小明：没什么意思。意思意思。领导：你这就不够意思了。小明：小意思，小意思。领导：你这人真有意思。小明：其实也没有别的意思。领导：那我就不好意思了。小明：是我不好意思。请问：以上“意思”分别是什么意思。
以上“意思”的含义如下：
1. 领导：你这是什么意思？ - 意思：指某种行为或言语所暗示的含义或意义。
2. 意思意思 - 意思：指某种行为或言语所暗示的含义或意义。
3. 小意思，小意思。 - 意思：指某种行为或言语所暗示的含义或意义，通常表示轻微或不重要的意思。
4. 其实也没有别的意思 - 意思：指某种行为或言语所暗示的含义或意义，表示实际上没有其他特定的含义或意义。
5. 是我不好意思 - 意思：指某种行为或言语所暗示的含义或意义，表示自己感到不好意思或尴尬。
```

千问的回答完全答非所问，不仅输出语言为英文，而且输出了 <|im\_start|> 等底层特殊控制符；开始强行解释 'that', 'meaning' 等毫无关联的英文单词。表现是灾难级的。

智谱的回答是相关的，浅层意义都成功解释了。但仔细阅读内容就会发现，它像一个只会查字典的机器人，完全屏蔽了对话中的社交潜台词（如：送礼、拒绝、客套），而是把每一个“意思”都翻译成了最安全的字面定义，甚至出现了“小意思就是不重要的意思”这种毫无信息量的同义反复（套娃）。表现只能说是及格。

### 3.总结分析(此部分对应作业要求 3：语言模型之间的横向对比分析)

通义千问-7B(Qwen-7B)是阿里云研发的通义千问大模型系列的 70 亿参数规模的模型。它是基于 Transformer 的大语言模型，在超大规模的预训练数据上进行训练得到。ChatGLM3 是智谱 AI 和清华大学 KEG 实验室联合发布的对话预训练模型。两者都是当代优秀的大语言模型。

在本次项目的五个地狱级测试问题中，两种大语言模型都展现了各自的生成倾向。对于偏向单一维度的题目(比如问题一和问题二)来说，两者各有胜负，通义千问 Qwen-7B-Chat 在生活常识的语境消歧上做出了较好的回答，但智谱 ChatGLM3-6B 则在复杂句法逻辑的反转拆解上表现得更为精准到位；对于逻辑嵌

套与分词较难的题目(比如问题三和问题四)来说,通义千问 Qwen-7B-Chat 虽然试图去推导你在问什么,但是最终注意力崩塌,给出了产生严重实体幻觉的错误结果,ChatGLM3-6B 则展现出对复杂推理的逃避,没有给出实质性的解析过程,而是给出捏造概念或原样照抄输入类似这样的敷衍性语句;对于难度极高的文化语境题目(比如问题五),通义千问 Qwen-7B-Chat 出现了底层指令解析的灾难性错误,直接逃逸到了毫无关联的英文输出,ChatGLM3-6B 的回答就比较没有贴合实际的人情语境,给出的回答太过于死板笼统,严重依赖字典式的废话模板套用,没有很好地理解所问问题中社交潜台词的关键所在。

两者都是优秀的大语言模型,经过这次深度的项目实验,我给出基于个人评测的意见:通义千问 Qwen-7B-Chat 具备良好的常识底座,但其注意力结构相对脆弱,遇到逻辑盲区时极易发生严重的实体幻觉和格式崩坏,未来需要在底层指令的鲁棒性上多加改进。ChatGLM3-6B 拥有极强的纯句法解析与排版能力,但在知识盲区时,更倾向于通过机械照抄、同义反复和伪造概念来维持表面上的“正确”,未来需要在人类真实社交语境的深度对齐上多加训练。